

Импорт библиотек

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.preprocessing import StandardScaler

from sklearn.model_selection import train_test_split

from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier,
GradientBoostingClassifier
from xgboost import XGBClassifier
from sklearn.metrics import accuracy_score, f1_score, roc_auc_score

from sklearn.model_selection import GridSearchCV

import numpy as np
```

1. Загрузка и предварительный анализ данных

```
data = pd.read_csv('WineQT.csv')
print(data.head())
print(data.info())
print(data.describe())
```

| | fixed acidity | volatile acidity | citric acid | residual sugar |
|---|---------------|------------------|-------------|----------------|
| 0 | 7.4 | 0.70 | 0.00 | 1.9 |
| 1 | 7.8 | 0.88 | 0.00 | 2.6 |
| 2 | 7.8 | 0.76 | 0.04 | 2.3 |
| 3 | 11.2 | 0.28 | 0.56 | 1.9 |
| 4 | 7.4 | 0.70 | 0.00 | 1.9 |

| | free sulfur dioxide | total sulfur dioxide | density | pH | sulphates |
|---|---------------------|----------------------|---------|------|-----------|
| 0 | 11.0 | 34.0 | 0.9978 | 3.51 | 0.56 |
| 1 | 25.0 | 67.0 | 0.9968 | 3.20 | 0.68 |
| 2 | 15.0 | 54.0 | 0.9970 | 3.26 | 0.65 |
| 3 | 17.0 | 60.0 | 0.9980 | 3.16 | 0.58 |

4 11.0 34.0 0.9978 3.51 0.56

```
alcohol  quality  Id
0      9.4      5    0
1      9.8      5    1
2      9.8      5    2
3      9.8      6    3
4      9.4      5    4
```

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 1143 entries, 0 to 1142

Data columns (total 13 columns):

| # | Column | Non-Null Count | Dtype |
|----|----------------------|----------------|---------|
| 0 | fixed acidity | 1143 non-null | float64 |
| 1 | volatile acidity | 1143 non-null | float64 |
| 2 | citric acid | 1143 non-null | float64 |
| 3 | residual sugar | 1143 non-null | float64 |
| 4 | chlorides | 1143 non-null | float64 |
| 5 | free sulfur dioxide | 1143 non-null | float64 |
| 6 | total sulfur dioxide | 1143 non-null | float64 |
| 7 | density | 1143 non-null | float64 |
| 8 | pH | 1143 non-null | float64 |
| 9 | sulphates | 1143 non-null | float64 |
| 10 | alcohol | 1143 non-null | float64 |
| 11 | quality | 1143 non-null | int64 |
| 12 | Id | 1143 non-null | int64 |

dtypes: float64(11), int64(2)

memory usage: 116.2 KB

None

| | fixed acidity | volatile acidity | citric acid | residual sugar \ |
|-------|---------------|------------------|-------------|------------------|
| count | 1143.000000 | 1143.000000 | 1143.000000 | 1143.000000 |
| mean | 8.311111 | 0.531339 | 0.268364 | 2.532152 |
| std | 1.747595 | 0.179633 | 0.196686 | 1.355917 |
| min | 4.600000 | 0.120000 | 0.000000 | 0.900000 |
| 25% | 7.100000 | 0.392500 | 0.090000 | 1.900000 |
| 50% | 7.900000 | 0.520000 | 0.250000 | 2.200000 |
| 75% | 9.100000 | 0.640000 | 0.420000 | 2.600000 |
| max | 15.900000 | 1.580000 | 1.000000 | 15.500000 |


| | chlorides | free sulfur dioxide | total sulfur dioxide |
|-----------|-------------|---------------------|----------------------|
| density \ | | | |
| count | 1143.000000 | 1143.000000 | 1143.000000 |
| mean | 0.086933 | 15.615486 | 45.914698 |
| std | 0.047267 | 10.250486 | 32.782130 |
| min | 0.012000 | 1.000000 | 6.000000 |

| | | | | | |
|----------|-------------|-------------|-------------|-------------|-------------|
| 25% | 0.070000 | 7.000000 | 21.000000 | | |
| 0.995570 | | | | | |
| 50% | 0.079000 | 13.000000 | 37.000000 | | |
| 0.996680 | | | | | |
| 75% | 0.090000 | 21.000000 | 61.000000 | | |
| 0.997845 | | | | | |
| max | 0.611000 | 68.000000 | 289.000000 | | |
| 1.003690 | | | | | |
| | pH | sulphates | alcohol | quality | Id |
| count | 1143.000000 | 1143.000000 | 1143.000000 | 1143.000000 | 1143.000000 |
| mean | 3.311015 | 0.657708 | 10.442111 | 5.657043 | 804.969379 |
| std | 0.156664 | 0.170399 | 1.082196 | 0.805824 | 463.997116 |
| min | 2.740000 | 0.330000 | 8.400000 | 3.000000 | 0.000000 |
| 25% | 3.205000 | 0.550000 | 9.500000 | 5.000000 | 411.000000 |
| 50% | 3.310000 | 0.620000 | 10.200000 | 6.000000 | 794.000000 |
| 75% | 3.400000 | 0.730000 | 11.100000 | 6.000000 | 1209.500000 |
| max | 4.010000 | 2.000000 | 14.900000 | 8.000000 | 1597.000000 |

2. Разведочный анализ данных (EDA)

Распределение целевой переменной:

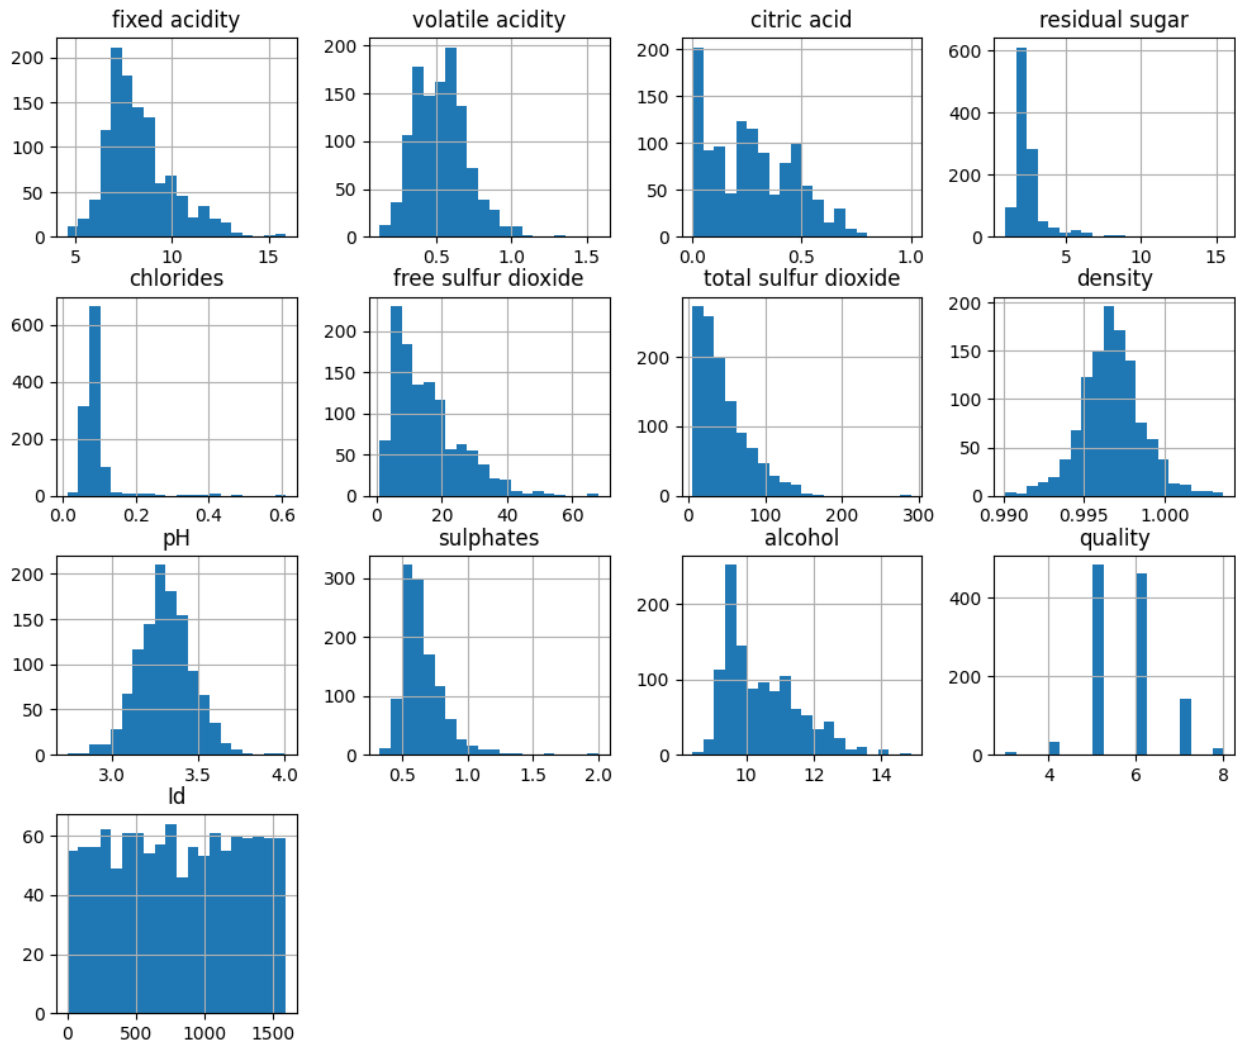
```
sns.countplot(x='quality', data=data)
plt.title('Распределение качества вина')
plt.show()
```



Наблюдается дисбаланс: большинство вин имеют оценку 5 или 6.

Гистограммы признаков:


```
data.hist(bins=20, figsize=(12, 10))  
plt.show()
```



Признаки имеют разный масштаб, что требует стандартизации.

Корреляционный анализ:

```
plt.figure(figsize=(10, 8))
sns.heatmap(data.corr(), annot=True, cmap='coolwarm')
plt.show()
```



Наибольшую корреляцию с quality имеют:

- alcohol (0.48)
- volatile acidity (-0.41)
- sulphates (0.26)

3. Подготовка данных

Кодирование целевой переменной (для классификации):

Преобразуем quality в 3 класса:

- Низкое (0): 0-4
- Среднее (1): 5-6
- Высокое (2): 7-10

```
data['quality_class'] = pd.cut(data['quality'], bins=[0, 4, 6, 10],
labels=[0, 1, 2])
```

```
X = data.drop(['quality', 'quality_class'], axis=1)
y = data['quality_class']
```

Масштабирование признаков:

```
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

4. Выбор метрик

Для задачи классификации с дисбалансом классов выберем:

- Accuracy (общая точность): одна из самых популярных метрик
- F1-score (среднее гармоническое precision и recall): подходит для задач с дисбалансом классов
- ROC-AUC (площадь под ROC-кривой, для многоклассовой классификации): для сравнения моделей на вероятностях, устойчива к дисбалансу

5. Разделение данных на train/test

```
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y,
test_size=0.2, random_state=42, stratify=y)
```

6. Выбор и обучение моделей (baseline)

Проверим 5 моделей:

- Логистическая регрессия (LogisticRegression)
- Метод опорных векторов (SVC)
- Случайный лес (RandomForestClassifier)
- Градиентный бустинг (GradientBoostingClassifier)
- XGBoost (XGBClassifier)

```
models = {
    'Logistic Regression': LogisticRegression(),
    'SVM': SVC(probability=True),
    'Random Forest': RandomForestClassifier(),
    'Gradient Boosting': GradientBoostingClassifier(),
    'XGBoost': XGBClassifier()
}

results = {}
for name, model in models.items():
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    results[name] = {
        'Accuracy': accuracy_score(y_test, y_pred),
        'F1-score': f1_score(y_test, y_pred, average='weighted'),
        'ROC-AUC': roc_auc_score(y_test, model.predict_proba(X_test),
multi_class='ovr')
```

```

    }

results_df = pd.DataFrame(results).T
print(results_df)

```

| | Accuracy | F1-score | ROC-AUC |
|---------------------|----------|----------|----------|
| Logistic Regression | 0.834061 | 0.808235 | 0.803709 |
| SVM | 0.851528 | 0.816625 | 0.795464 |
| Random Forest | 0.903930 | 0.884047 | 0.842532 |
| Gradient Boosting | 0.877729 | 0.862444 | 0.775337 |
| XGBoost | 0.886463 | 0.870817 | 0.787009 |

7. Подбор гиперпараметров

```

# Определим параметры для каждой модели
param_grids = {
    'Logistic Regression': {
        'C': [0.1, 1],
        'penalty': ['l2']
    },
    'SVM': {
        'C': [0.1, 1, 10],
        'kernel': ['linear', 'rbf', 'poly'],
        'gamma': ['scale', 'auto']
    },
    'Random Forest': {
        'n_estimators': [75, 80, 85, 90, 95, 100, 105, 110, 115, 120],
        'max_depth': [10, 15, 20, 25]
    },
    'Gradient Boosting': {
        'n_estimators': [50, 100, 200],
        'learning_rate': [0.01, 0.1, 0.2],
        'max_depth': [3, 5, 7]
    },
    'XGBoost': {
        'n_estimators': [50, 100, 200],
        'learning_rate': [0.01, 0.1, 0.2],
        'max_depth': [3, 5, 7],
        'subsample': [0.8, 1.0],
        'colsample_bytree': [0.8, 1.0]
    }
}

best_models = {}
for name, model in models.items():
    print(f"Подбор параметров для модели: {name}")
    grid_search = GridSearchCV(
        estimator=model,
        param_grid=param_grids[name],
        cv=5,

```



```

        scoring='accuracy',
        n_jobs=-1,
        verbose=1
    )
    grid_search.fit(X_train, y_train)
    best_models[name] = grid_search.best_estimator_
    print(f"Лучшие параметры для {name}: {grid_search.best_params_}\n")

```

Подбор параметров для модели: Logistic Regression

Fitting 5 folds for each of 2 candidates, totalling 10 fits

Лучшие параметры для Logistic Regression: {'C': 0.1, 'penalty': 'l2'}

Подбор параметров для модели: SVM

Fitting 5 folds for each of 18 candidates, totalling 90 fits

Лучшие параметры для SVM: {'C': 10, 'gamma': 'auto', 'kernel': 'rbf'}

Подбор параметров для модели: Random Forest

Fitting 5 folds for each of 40 candidates, totalling 200 fits

Лучшие параметры для Random Forest: {'max_depth': 20, 'n_estimators': 105}

Подбор параметров для модели: Gradient Boosting

Fitting 5 folds for each of 27 candidates, totalling 135 fits

Лучшие параметры для Gradient Boosting: {'learning_rate': 0.1, 'max_depth': 7, 'n_estimators': 100}

Подбор параметров для модели: XGBoost

Fitting 5 folds for each of 108 candidates, totalling 540 fits

Лучшие параметры для XGBoost: {'colsample_bytree': 0.8, 'learning_rate': 0.2, 'max_depth': 3, 'n_estimators': 200, 'subsample': 1.0}

8. Сравнение моделей с оптимальными гиперпараметрами с baseline

```

optimized_results = {}
for name, model in best_models.items():
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    optimized_results[name] = {
        'Accuracy': accuracy_score(y_test, y_pred),
        'F1-score': f1_score(y_test, y_pred, average='weighted'),
        'ROC-AUC': roc_auc_score(y_test, model.predict_proba(X_test),
multi_class='ovr')
    }

optimized_results_df = pd.DataFrame(optimized_results).T
print("\nРезультаты после подбора гиперпараметров:")
print(optimized_results_df)

```

```
# Сравним с baseline
print("\nСравнение с baseline:")
comparison_df = pd.concat([results_df, optimized_results_df],
keys=['Baseline', 'Optimized'], axis=1)
print(comparison_df)
```

Результаты после подбора гиперпараметров:

| | Accuracy | F1-score | ROC-AUC |
|---------------------|----------|----------|----------|
| Logistic Regression | 0.838428 | 0.809073 | 0.792483 |
| SVM | 0.877729 | 0.861198 | 0.798737 |
| Random Forest | 0.908297 | 0.889094 | 0.864533 |
| Gradient Boosting | 0.899563 | 0.882744 | 0.803974 |
| XGBoost | 0.895197 | 0.878750 | 0.759518 |

Сравнение с baseline:

| | Baseline | | | Optimized | |
|---------------------|----------|----------|----------|-----------|----------|
| \ | Accuracy | F1-score | ROC-AUC | Accuracy | F1-score |
| Logistic Regression | 0.834061 | 0.808235 | 0.803709 | 0.838428 | 0.809073 |
| SVM | 0.851528 | 0.816625 | 0.795464 | 0.877729 | 0.861198 |
| Random Forest | 0.903930 | 0.884047 | 0.842532 | 0.908297 | 0.889094 |
| Gradient Boosting | 0.877729 | 0.862444 | 0.775337 | 0.899563 | 0.882744 |
| XGBoost | 0.886463 | 0.870817 | 0.787009 | 0.895197 | 0.878750 |

| | ROC-AUC |
|---------------------|----------|
| Logistic Regression | 0.792483 |
| SVM | 0.798737 |
| Random Forest | 0.864533 |
| Gradient Boosting | 0.803974 |
| XGBoost | 0.759518 |

Exception ignored in: <function ResourceTracker.__del__ at 0x107edfb00>

Traceback (most recent call last):

```
File
"/opt/homebrew/Cellar/python@3.13/3.13.3/Frameworks/Python.framework/
Versions/3.13/lib/python3.13/multiprocessing/resource_tracker.py",
line 82, in __del__
File
"/opt/homebrew/Cellar/python@3.13/3.13.3/Frameworks/Python.framework/
Versions/3.13/lib/python3.13/multiprocessing/resource_tracker.py",
line 91, in _stop
```

```
File
"/opt/homebrew/Cellar/python@3.13/3.13.3/Frameworks/Python.framework/
Versions/3.13/lib/python3.13/multiprocessing/resource_tracker.py",
line 116, in _stop_locked
ChildProcessError: [Errno 10] No child processes
Exception ignored in: <function ResourceTracker.__del__ at
0x104fe3b00>
Traceback (most recent call last):
  File
"/opt/homebrew/Cellar/python@3.13/3.13.3/Frameworks/Python.framework/
Versions/3.13/lib/python3.13/multiprocessing/resource_tracker.py",
line 82, in __del__
  File
"/opt/homebrew/Cellar/python@3.13/3.13.3/Frameworks/Python.framework/
Versions/3.13/lib/python3.13/multiprocessing/resource_tracker.py",
line 91, in _stop
  File
"/opt/homebrew/Cellar/python@3.13/3.13.3/Frameworks/Python.framework/
Versions/3.13/lib/python3.13/multiprocessing/resource_tracker.py",
line 116, in _stop_locked
ChildProcessError: [Errno 10] No child processes
Exception ignored in: <function ResourceTracker.__del__ at
0x1054cbb00>
Traceback (most recent call last):
  File
"/opt/homebrew/Cellar/python@3.13/3.13.3/Frameworks/Python.framework/
Versions/3.13/lib/python3.13/multiprocessing/resource_tracker.py",
line 82, in __del__
  File
"/opt/homebrew/Cellar/python@3.13/3.13.3/Frameworks/Python.framework/
Versions/3.13/lib/python3.13/multiprocessing/resource_tracker.py",
line 91, in _stop
  File
"/opt/homebrew/Cellar/python@3.13/3.13.3/Frameworks/Python.framework/
Versions/3.13/lib/python3.13/multiprocessing/resource_tracker.py",
line 116, in _stop_locked
ChildProcessError: [Errno 10] No child processes
Exception ignored in: <function ResourceTracker.__del__ at
0x106d97b00>
Traceback (most recent call last):
  File
"/opt/homebrew/Cellar/python@3.13/3.13.3/Frameworks/Python.framework/
Versions/3.13/lib/python3.13/multiprocessing/resource_tracker.py",
line 82, in __del__
  File
"/opt/homebrew/Cellar/python@3.13/3.13.3/Frameworks/Python.framework/
Versions/3.13/lib/python3.13/multiprocessing/resource_tracker.py",
line 91, in _stop
  File
```

```
"/opt/homebrew/Cellar/python@3.13/3.13.3/Frameworks/Python.framework/
Versions/3.13/lib/python3.13/multiprocessing/resource_tracker.py",
line 116, in _stop_locked
ChildProcessError: [Errno 10] No child processes
Exception ignored in: <function ResourceTracker.__del__ at
0x10505bb00>
Traceback (most recent call last):
  File
"/opt/homebrew/Cellar/python@3.13/3.13.3/Frameworks/Python.framework/
Versions/3.13/lib/python3.13/multiprocessing/resource_tracker.py",
line 82, in __del__
  File
"/opt/homebrew/Cellar/python@3.13/3.13.3/Frameworks/Python.framework/
Versions/3.13/lib/python3.13/multiprocessing/resource_tracker.py",
line 91, in _stop
  File
"/opt/homebrew/Cellar/python@3.13/3.13.3/Frameworks/Python.framework/
Versions/3.13/lib/python3.13/multiprocessing/resource_tracker.py",
line 116, in _stop_locked
ChildProcessError: [Errno 10] No child processes
Exception ignored in: <function ResourceTracker.__del__ at
0x1027dfb00>
Traceback (most recent call last):
  File
"/opt/homebrew/Cellar/python@3.13/3.13.3/Frameworks/Python.framework/
Versions/3.13/lib/python3.13/multiprocessing/resource_tracker.py",
line 82, in __del__
  File
"/opt/homebrew/Cellar/python@3.13/3.13.3/Frameworks/Python.framework/
Versions/3.13/lib/python3.13/multiprocessing/resource_tracker.py",
line 91, in _stop
  File
"/opt/homebrew/Cellar/python@3.13/3.13.3/Frameworks/Python.framework/
Versions/3.13/lib/python3.13/multiprocessing/resource_tracker.py",
line 116, in _stop_locked
ChildProcessError: [Errno 10] No child processes
Exception ignored in: <function ResourceTracker.__del__ at
0x104593b00>
Traceback (most recent call last):
  File
"/opt/homebrew/Cellar/python@3.13/3.13.3/Frameworks/Python.framework/
Versions/3.13/lib/python3.13/multiprocessing/resource_tracker.py",
line 82, in __del__
  File
"/opt/homebrew/Cellar/python@3.13/3.13.3/Frameworks/Python.framework/
Versions/3.13/lib/python3.13/multiprocessing/resource_tracker.py",
line 91, in _stop
  File
"/opt/homebrew/Cellar/python@3.13/3.13.3/Frameworks/Python.framework/
```

```

Versions/3.13/lib/python3.13/multiprocessing/resource_tracker.py",
line 116, in _stop_locked
ChildProcessError: [Errno 10] No child processes
Exception ignored in: <function ResourceTracker.__del__ at
0x1046bbb00>
Traceback (most recent call last):
  File
"/opt/homebrew/Cellar/python@3.13/3.13.3/Frameworks/Python.framework/
Versions/3.13/lib/python3.13/multiprocessing/resource_tracker.py",
line 82, in __del__
  File
"/opt/homebrew/Cellar/python@3.13/3.13.3/Frameworks/Python.framework/
Versions/3.13/lib/python3.13/multiprocessing/resource_tracker.py",
line 91, in _stop
  File
"/opt/homebrew/Cellar/python@3.13/3.13.3/Frameworks/Python.framework/
Versions/3.13/lib/python3.13/multiprocessing/resource_tracker.py",
line 116, in _stop_locked
ChildProcessError: [Errno 10] No child processes

```

9. Выводы

- Лучшая модель: Random Forest (F1-score = 0.889).
- Ансамблевые методы (Random Forest, Gradient Boosting, XGBoost) показали себя лучше других.
- Logistic Regression слабо реагирует на подбор параметров

Визуализация сравнения Baseline и Optimized моделей

```

metrics = ['Accuracy', 'F1-score', 'ROC-AUC']
n_metrics = len(metrics)
n_models = len(results_df)

plt.figure(figsize=(15, 5 * n_metrics))
for i, metric in enumerate(metrics, 1):
    plt.subplot(n_metrics, 1, i)
    baseline_values = results_df[metric]
    optimized_values = optimized_results_df[metric]

    x = range(n_models)
    width = 0.35


    plt.bar(x, baseline_values, width, label='Baseline',
            color='skyblue')
    plt.bar([p + width for p in x], optimized_values, width,
            label='Optimized', color='orange')

    plt.xticks([p + width/2 for p in x], results_df.index)
    plt.ylabel(metric)
    plt.title(f'Сравнение {metric} до и после оптимизации')

```

```
plt.legend()
plt.grid(True, linestyle='--', alpha=0.5)

plt.tight_layout()
plt.show()
```



Влияние гиперпараметров на качество (для Random Forest)

```
param1 = 'n_estimators'
param2 = 'max_depth'

grid_values = param_grids['Random Forest']
grid_search = GridSearchCV(RandomForestClassifier(),
```

```

param_grids['Random Forest'], cv=5, scoring='accuracy')
grid_search.fit(X_train, y_train)

results = grid_search.cv_results_
scores =
np.array(results['mean_test_score']).reshape(len(grid_values[param1]),
len(grid_values[param2]))

plt.figure(figsize=(10, 6))
sns.heatmap(scores, annot=True, fmt=".3f",
            xticklabels=grid_values[param2],
            yticklabels=grid_values[param1],
            cmap="viridis")
plt.xlabel(param2)
plt.ylabel(param1)
plt.title('Тепловая карта зависимости Accuracy от гиперпараметров
(Random Forest)')
plt.show()

```

